

The Programming Interface

Instructor: Youngjin Kwon

Let's think about interface

- OS designer
 - Know internals of how OS works
 - Know internals of how hardware works
- Application developer
 - Do not want to know internals of OS and hardware
- OS provides a simplistic view to use and manage OS internal resources by **APIs**
 - System calls

What to consider designing APIs?

- **What functionality should be in kernel?**
 - What functionality should be in user library?
- How should OS be organized?

Some thoughts

- Let's focus on building APIs for process creation
- We can put the functionality of managing processes and IO in user-level
 - *What is a problem?*


Some thoughts (Cont'd)

- **So, the functionality must be in the kernel**
- Then, how to design the APIs?
 - Designing good APIs is a hard problem
 - Our approach: Let's look at use cases and build APIs to reflect the use cases


Main Points

- Creating and managing processes
 - fork, exec, wait
- Performing I/O
 - open, read, write, close
- Monolithic kernel & Microkernel

Use case: Shell



- A shell is a job control system
 - Allows programmer to create and manage a set of programs to do some tasks
 - Windows, MacOS, Linux all have shells



- Example: to compile a C program

`$ cc -c sourcefile1.c` ➡ It requires create a process and write data to sourcefile1.o

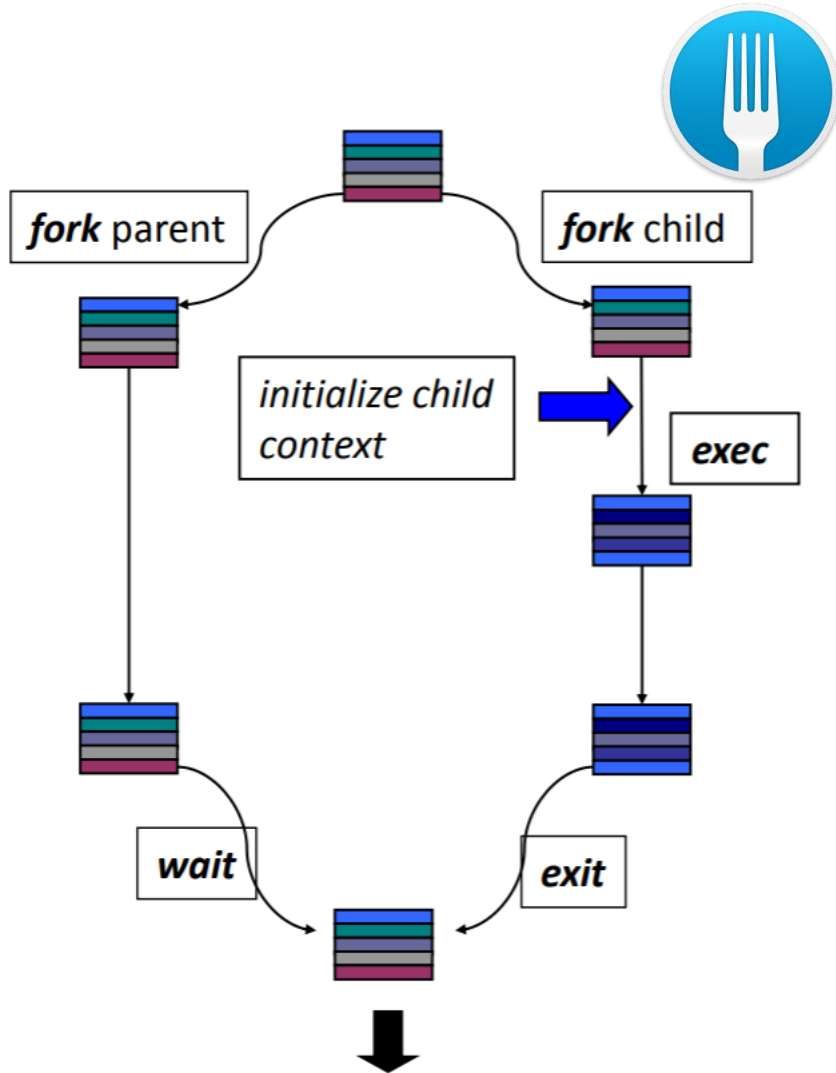
`$ cc -c sourcefile2.c`

`$ ln -o program sourcefile1.o sourcefile2.o`

Design questions

- If the shell runs at user-level, what system calls does it make to run each of the programs?
 - Ex: gcc, ln
- Why does process creation need system calls?
 - Why kernel-level?
 - Who has to manage process states? A parent process or kernel?

UNIX Process Management APIs



```
int pid = fork();
```

Create a new process that is a clone of its parent.

```
exec*("program" [, argvp, envp]);
```

Overlay the calling process virtual memory with a new program, and transfer control to it.

```
exit(status);
```

Exit with status, destroying the process.

```
int pid = wait*(&status);
```

Wait for exit (or other status change) of a child.

Question: What does this code print?

```
int child_pid = fork();  
if (child_pid == 0) {           // I'm the child process  
    printf("I am process # %d\n", getpid());  
    return 0;  
} else {                       // I'm the parent process  
    printf("I am parent of process # %d\n", child_pid);  
    return 0;  
}
```

Fork() and exec()

```
pid_t pid = 0; int status;
```

```
pid = fork();
```

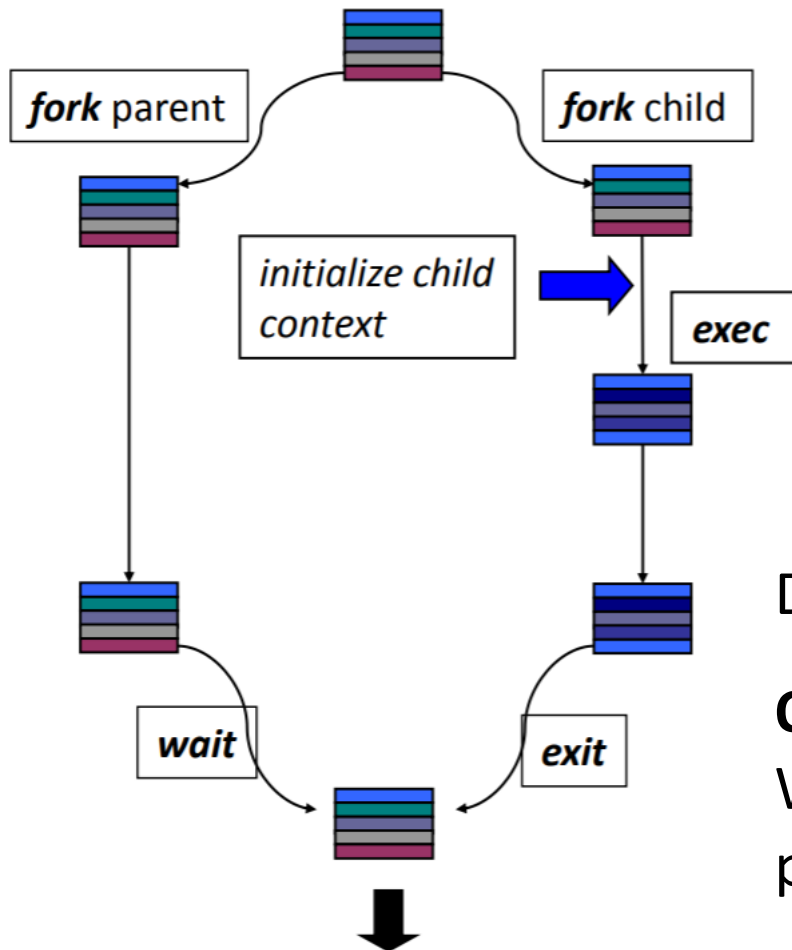
```
if (pid == 0) {  
    printf("I am the child.");  
    execl("/bin/ls", "ls", "-l", "/usr/lib", (char *) 0);  
    // not reachable  
    perror("In exec(): ");  
}
```

```
if (pid > 0) {  
    printf("I am the parent, and the child is %d.\n", pid);  
    pid = wait(&status);  
    printf("End of process %d: ", pid);  
}
```

Questions

- Can UNIX `fork()` return an error? Why?
- Can UNIX `exec()` return an error? Why?
- Can UNIX `wait()` ever return immediately? Why?

UNIX Process Management APIs



```
$ cc -c sourcefile1.c
```

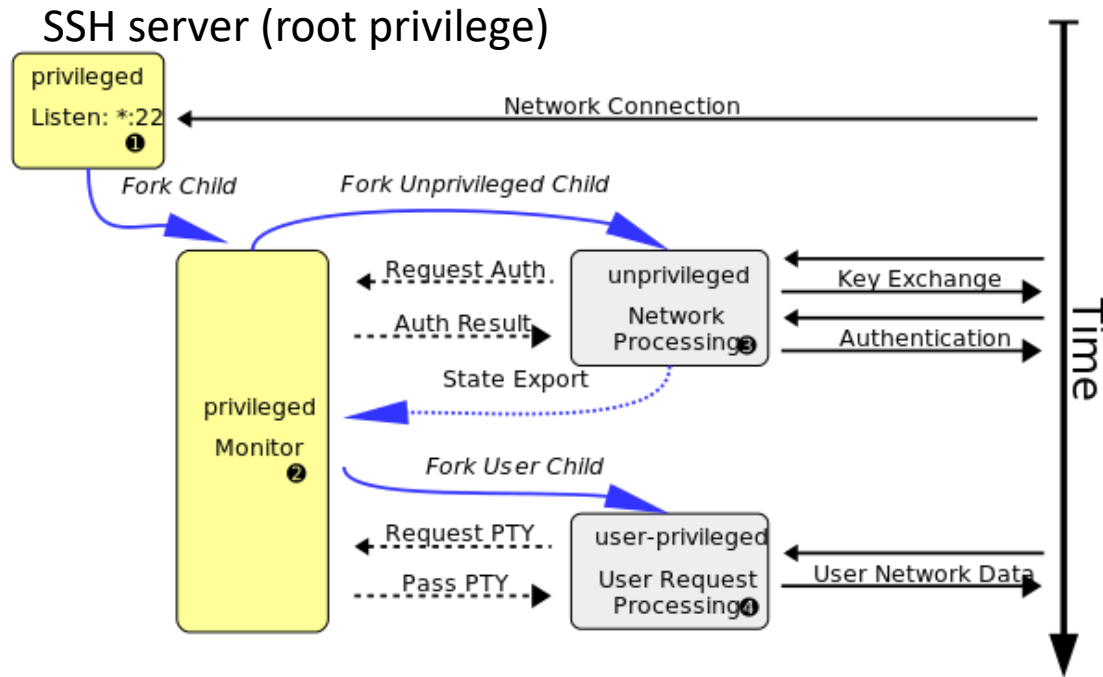
Bash (\$) → Bash (child)
→ Run compiler (cc)
→ [compile]
→ exit
→ return to Bash (\$)

Does execution path above match APIs?

Question: Why fork?

What if just run compile in the same process as bash (\$)?

SSH: Privilege separation



Question: Why fork?

For separating privilege

1. SSH server must be protected from malicious clients.
2. SSH shell should be run as a user privilege, not root.

Some thoughts (Cont'd)

- So, the functionality must be in the kernel
- **Then, how to design the APIs?**
 - Designing good APIs is a hard problem
 - **Our approach: Let's look at use cases and build APIs to reflect the use cases**

John Ousterhout's google talk

- https://www.youtube.com/watch?v=bmSAYlu0NcY&ab_channel=TalksatGoogle
- If you have to pick a most important idea in CS, what would you pick? (3:00)
- Professors are not eligible to teach code skills ☹️ (5:40)
- What is good interfaces? Deep interface (13:00)
- Example of Deep interface (20:32)

UNIX I/O API model (POSIX model)

- **Uniformity** (*Unix treats everything as a file*)
 - All operations on all files, devices use the same set of system calls: open, close, read, write
- Open before use
 - Open returns a handle (file descriptor) for use in later calls on the file
- Byte-oriented (not block-oriented)
- Kernel-buffered read/write
- Explicit close
 - To garbage collect the open file descriptor

Implementing UNIX fork

Steps to implement UNIX fork

- Create and initialize the process control block (PCB) in the kernel
- Create a new address space
- *Initialize the address space with a copy of the entire contents of the address space of the parent*
- Inherit the execution context of the parent (e.g., any open files)
- Inform the scheduler that the new process is ready to run

Implementing UNIX exec

- Steps to implement UNIX exec
 - Load the program into the current address space
 - Copy arguments into memory in the address space
 - Initialize the hardware context to start execution at ``start''

Implementing a Shell

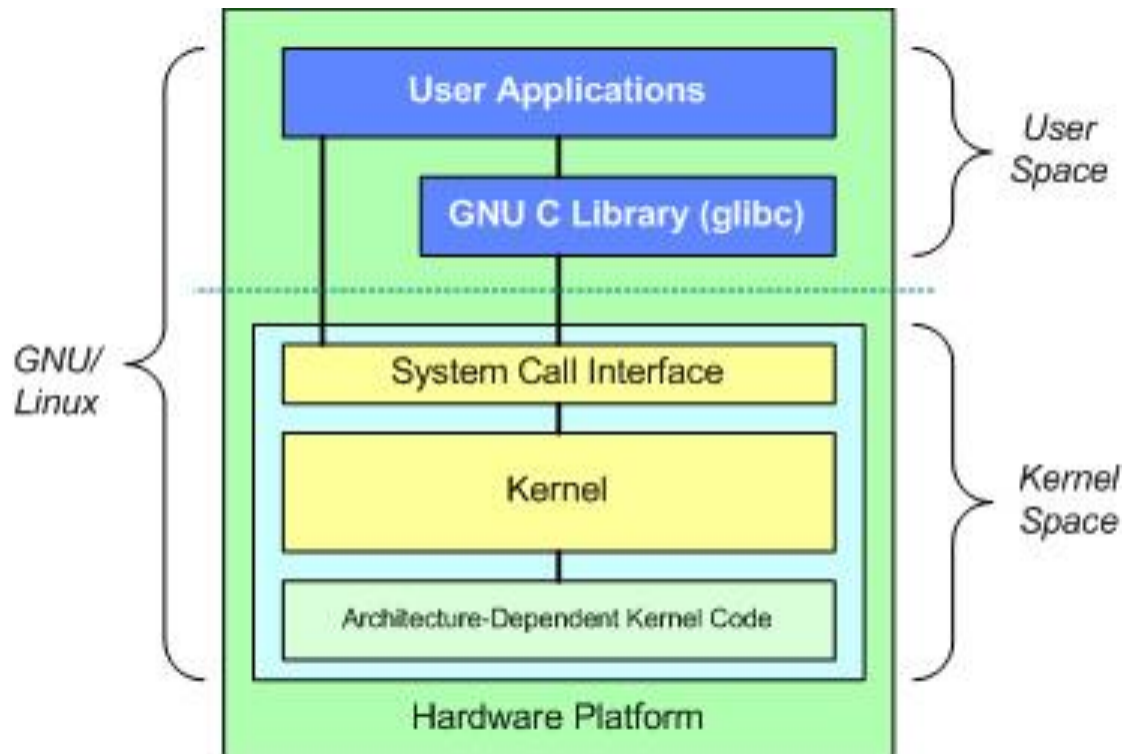
```
char *prog, **args;
int child_pid;

// Read and parse the input a line at a time
while (readAndParseCmdLine(&prog, &args)) {
    child_pid = fork();    // create a child process
    if (child_pid == 0) {
        exec(prog, args);    // I'm the child process. Run program
        // NOT REACHED
    } else {
        wait(child_pid);    // I'm the parent, wait for child
        return 0;
    }
}
```

What to consider designing APIs?

- What functionality should be in kernel?
 - **What functionality should be in user library?**
- How should OS be organized?

Standard C library



Why is the standard library required?

A system call allows a program to request a service—for example, open a file or create a new process—from the kernel. At the assembler level, making a system call requires the caller to assign the unique system call number and the argument values to particular registers, and then execute a special instruction (e.g., SYSENTER on modern x86 architectures) that switches the processor to kernel mode to execute the system-call handling code. Upon return, the kernel places the system call's result status into a particular register and executes a special instruction (e.g., SYSEXIT on x86) that returns the processor to user mode. The usual convention for the result status is that a non-negative value means success, while a negative value means failure. A negative result status is the negated error number (`errno`) that indicates the cause of the failure.

All of the details of making a system call are normally hidden from the user by the C library, which provides a corresponding wrapper function and header file definitions for most system calls. The wrapper function accepts the system call arguments as function

<https://lwn.net/Articles/534682/>

What to consider designing APIs?

- What functionality should be in kernel?
- What functionality should be in user library?


→ **How should OS be organized?**

option 1: All functionalities in a single kernel

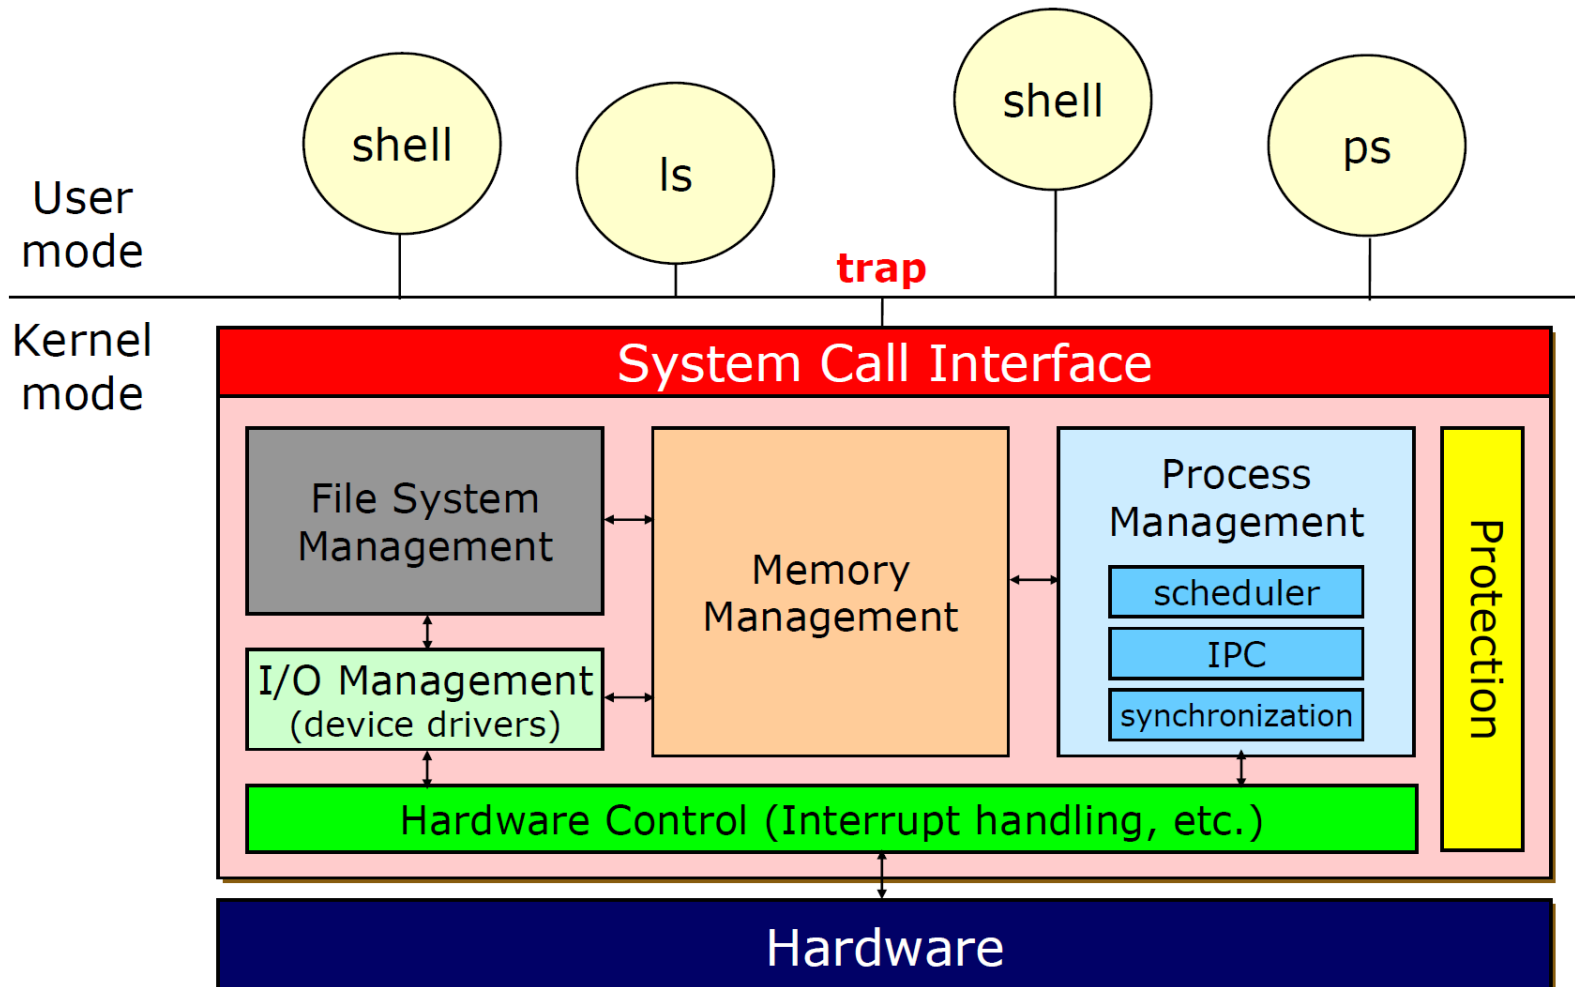
option 2: Functionalities are partitioned to kernel and several services

Monolithic Kernel

- All OS functionality is included in the kernel
 - Applications do not have control over resources
- Advantages
 - well-understood, **good performance**
- Disadvantages
 - **No protection btw kernel components**
 - Not easily extended/modified
- Example:
 - Windows, BSD, Linux




Operating System Structure



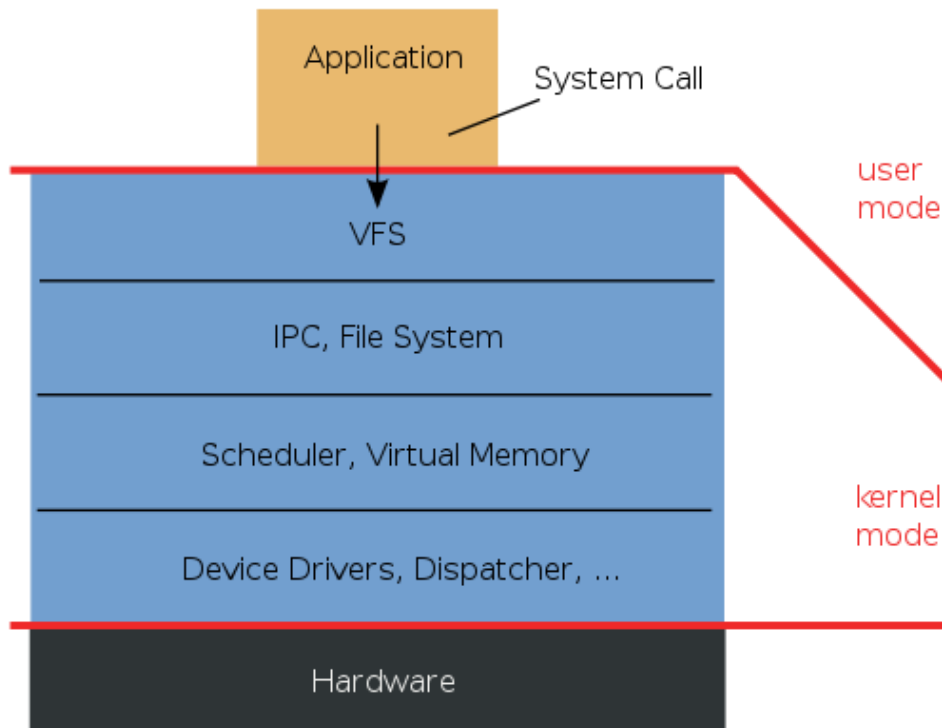
Microkernel

- OS kernel is very small – minimal functionality
 - The microkernel layer provides a set of minimal core services and is the interface to the hardware layer
- Other services (drivers, memory managers, etc.) are implemented as separate modules at user level
- Examples: Mach, L4

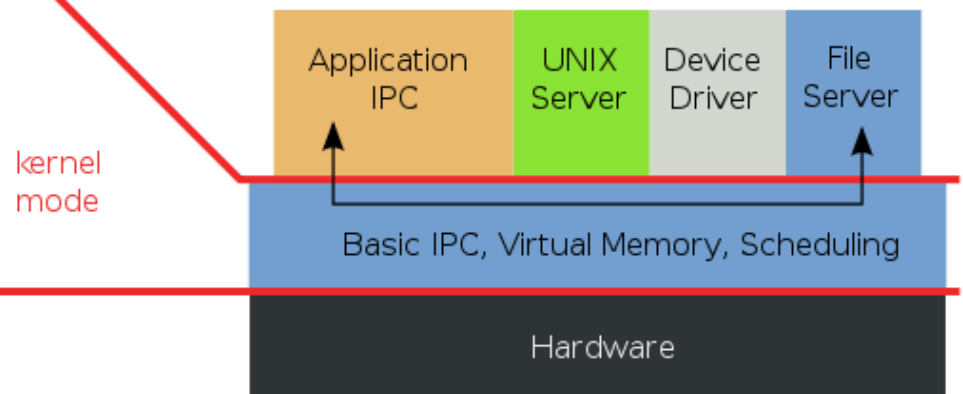


Monolithic Kernels VS Microkernels

Monolithic Kernel based Operating System



Microkernel based Operating System



Microkernel - Advantages

- Modularity
- Flexibility and extensibility
 - Easier to replace modules – fewer dependencies
 - Different servers can implement the same service in different ways
- **Safety (protection)**
 - Each server is protected by the OS from other servers
 - Servers are isolated; errors in one don't affect others
- Correctness
 - Easier to verify a small kernel

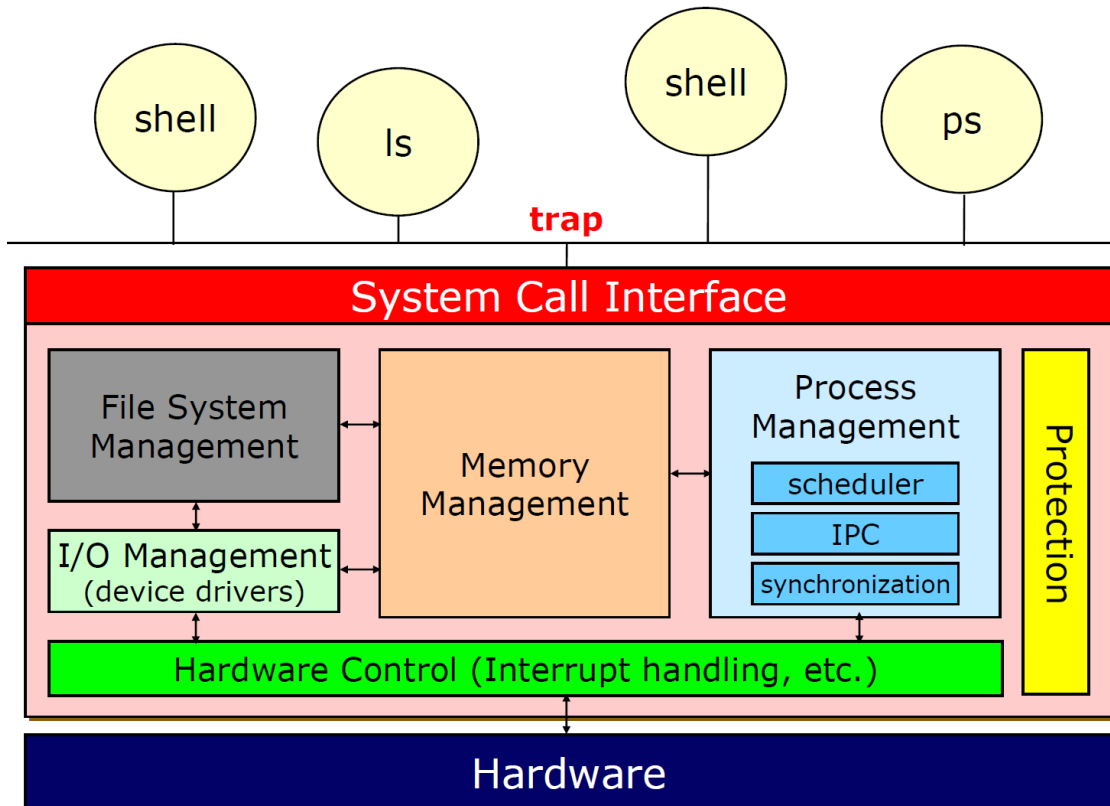
Microkernels- Disadvantages

- Performance issue
 - Slow – due to “cross-domain” information transfers?
 - Server-to-OS, OS-to-server IPC is thought to be a major source of inefficiency
 - Much faster to communicate between two modules that are both in OS

Hierarchy of Communication Mechanisms

- Fastest to Slowest
 - Function calls within same process
 - System calls (mode switch)
 - Context switch (process switch)
 - IPC between processes (message passing, on the same or different machines)

Big picture: where are we?



We've discussed:

- Protection
- Hardware control
- Scheduling
- APIs for process management
- APIs for IO